

Simulation and Empirical Analysis of Linear Regression Using Multivariate Normal, FirstYearGPA, Prostate, and Job Proficiency Data

```
chooseCRANmirror(graphics = FALSE, ind = 68)

#QUESTION #1:
library(MASS)

simulate_regression <- function(Sigma, n = 200, reps = 3) {
  results <- list()

  for (i in 1:reps) {
    X <- mvrnorm(n, mu = rep(0, 4), Sigma = Sigma)
    colnames(X) <- c("X1", "X2", "X3", "X4")
    dat <- as.data.frame(X)

    dat$Y <- 0 + 1*X[,1]

    model <- lm(Y ~ X1)

    vif_values <- sapply(
      1 / (1 - summary(coef(model))^2)
    )
    names(vif_values) <- colnames(X)

    results[[i]] <- list(
      coefficients = summary(model)$coefficients,
      vif = vif_values,
      cor = cor(X)
    )
  }

  return(results)
}

Sigma_indep <- diag(4)
Sigma_corr <- matrix(c(
  1, 0, 0.95, 0,
  0, 1, 0, -0.95,
  0.95, 0, 1, 0,
  0, -0.95, 0, 1
), nrow = 4)

#Case 1:
set.seed(123)
results_indep <- simulate_regression(Sigma_indep)
```

```
cat("### First Replication Results (Independent Predictors)\n")
```

```
## ### First Replication Results (Independent Predictors)
```

```
print(results_indep[[1]]$coefficients)
```

```
##           Estimate Std. Error   t value   Pr(>|t|)
## (Intercept)  0.03120373 0.07338407  0.4252113 6.711515e-01
## X1           0.92846990 0.07200968 12.8936809 6.528325e-28
## X2           0.97631222 0.07629346 12.7968014 1.286124e-27
## X3           0.11134758 0.07410124  1.5026412 1.345493e-01
## X4          -0.05739540 0.07799656 -0.7358710 4.626935e-01
```

```
cat("\nVIF:\n")
```

```
##
## VIF:
```

```
print(results_indep[[1]]$vif)
```

```
##           X1           X2           X3           X4
## 1.009680 1.004884 1.010096 1.003653
```

```
cat("\nSample Correlation Matrix:\n")
```

```
##
## Sample Correlation Matrix:
```

```
print(results_indep[[1]]$cor)
```

```
##           X1           X2           X3           X4
## X1 1.00000000 0.03197187 0.08015345 -0.04640932
## X2 0.03197187 1.00000000 -0.05141456 -0.03023254
## X3 0.08015345 -0.05141456 1.00000000 -0.02770462
## X4 -0.04640932 -0.03023254 -0.02770462 1.00000000
```

```
#Case 2:
```

```
set.seed(123)
```

```
results_corr <- simulate_regression(Sigma_corr)
```

```
cat("### First Replication Results (Correlated Predictors)\n")
```

```
## ### First Replication Results (Correlated Predictors)
```

```
print(results_corr[[1]]$coefficients)
```

```
##           Estimate Std. Error  t value    Pr(>|t|)
## (Intercept) 0.03120373 0.07338407 0.4252113 6.711515e-01
## X1          0.75850820 0.23029226 3.2936766 1.174282e-03
## X2          0.99202650 0.24161051 4.1058913 5.920384e-05
## X3          0.22232358 0.23200523 0.9582697 3.391134e-01
## X4          0.11743570 0.24661864 0.4761834 6.344768e-01
```

```
cat("\nVIF:\n")
```

```
##
## VIF:
```

```
print(results_corr[[1]]$vif)
```

```
##           X1           X2           X3           X4
## 8.940914 10.844498 8.969170 10.814124
```

```
cat("\nSample Correlation Matrix:\n")
```

```
##
## Sample Correlation Matrix:
```

```
print(results_corr[[1]]$cor)
```

```
##           X1           X2           X3           X4
## X1 1.00000000 0.03771533 0.94217707 -0.02396935
## X2 0.03771533 1.00000000 0.06181773 -0.95251652
## X3 0.94217707 0.06181773 1.00000000 -0.04433179
## X4 -0.02396935 -0.95251652 -0.04433179 1.00000000
```

```
#Case 3:
```

```
cat("### Standard Errors Across Replications (Independent vs Correlated)\n")
```

```
## ### Standard Errors Across Replications (Independent vs Correlated)
```

```
se_comparison <- data.frame(
  Replication = rep(1:3, each = 2),
  Case = rep(c("Independent", "Correlated"), 3),
  SE_X1 = c(
    results_indep[[1]]$coefficients["X1", "Std. Error"],
    results_corr[[1]]$coefficients["X1", "Std. Error"],
    results_indep[[2]]$coefficients["X1", "Std. Error"],
    results_corr[[2]]$coefficients["X1", "Std. Error"],
    results_indep[[3]]$coefficients["X1", "Std. Error"],
    results_corr[[3]]$coefficients["X1", "Std. Error"]
  )
)

print(se_comparison)
```

##	Replication	Case	SE_X1
## 1	1	Independent	0.07200968
## 2	1	Correlated	0.23029226
## 3	2	Independent	0.07241033
## 4	2	Correlated	0.24743229
## 5	3	Independent	0.07643216
## 6	3	Correlated	0.24169974

#Comments: In the independent predictors case, the standard errors for all #coefficients were small (around 0.07-0.08), leading to highly significant #t-tests for x1 and x2 ($p < 0.0001$) and non-significant results for x3 and #x4 ($p > 0.1$), correctly reflecting their true beta values of 1, 1, 0, and 0 #respectively. The correlated predictors case showed substantially inflated #standard errors (0.23-0.25), about 3-4 times larger than the independent case. #This inflation reduced the significance of the t-tests, while x1 and x2 #remained significant ($p = 0.001$ and $p < 0.0001$ respectively), their test #statistics were much weaker ($t = 3.3$ vs 12.9 for x1, $t = 4.1$ vs 12.8 for x2).

#The dramatic difference in outcomes stems from multicollinearity. The #correlation matrices reveal near-perfect correlations between x1-x3 (0.94) and #x2-x4 (-0.95) in the correlated case, while the independent case showed #near-zero correlations. This is reflected in the VIF values: all VIFs were #around 1 (no multicollinearity) in the independent case, but ranged from #8.9-10.8 in the correlated case, indicating severe multicollinearity. When #predictors are highly correlated, it becomes difficult for the model to isolate #their individual effects, inflating the standard errors of their coefficient #estimates.

#Across three replications (as shown in the SE comparison table), the pattern #remained consistent: Independent case maintained small, stable standard errors #(0.07-0.08 for x1). Correlated case consistently showed 3-4 times larger #standard errors (0.22-0.25 for x1). While the exact p-values varied slightly #across replications, the fundamental pattern held: x1 and x2 were always highly #significant in independent case, x1 and x2 remained significant but with weaker #evidence in correlated case, and x3 and x4 were never significant in either #case.

#Multicollinearity creates a shared variance problem among predictors. When x1 #and x3 are highly correlated ($r = 0.94$), the model struggles to determine which #predictor is truly responsible for explaining variation in Y. This uncertainty #shows up as inflated standard errors, we become less confident about each #predictor's unique contribution. The effect is particularly noticeable for x1 #(true beta=1), whose significance was dramatically reduced, because its effect #could be partially explained by its correlated counterpart x3 (true beta=0). #The VIF values quantify this inflation, showing the variance of each #coefficient estimate is inflated 8-10 times due to these correlations. While #the model can still detect the true effects (x1 and x2 remain significant), it #does so with much less precision.

#QUESTION #2:

`install.packages("Stat2Data")`

##

The downloaded binary packages are in

```
## /var/folders/n2/q439mgrn0rqd6xgjzjkj652r0000gn/T//RtmptocgW8/downloaded_packages
```

```
library(Stat2Data)
data("FirstYearGPA")

gpa_data <- FirstYearGPA[, c("GPA", "HSGPA", "SATV", "SATM", "HU", "White")]

full_model <- lm(GPA ~ HSGPA + SATV + SATM + HU + White +
                 HSGPA:White + SATV:White + SATM:White + HU:White,
                 data = gpa_data)

reduced_model <- lm(GPA ~ HSGPA + SATV + SATM + HU + White,
                   data = gpa_data)

anova_result <- anova(reduced_model, full_model)

cat("F-test results comparing models:\n")
```

```
## F-test results comparing models:
```

```
print(anova_result)
```

```
## Analysis of Variance Table
```

```
##
```

```
## Model 1: GPA ~ HSGPA + SATV + SATM + HU + White
```

```
## Model 2: GPA ~ HSGPA + SATV + SATM + HU + White + HSGPA:White + SATV:White +
```

```
## SATM:White + HU:White
```

```
## Res.Df RSS Df Sum of Sq F Pr(>F)
```

```
## 1 213 31.265
```

```
## 2 209 30.429 4 0.83609 1.4356 0.2234
```

```
p_value <- anova_result$`Pr(>F)`[2]
cat("\nThe p-value for the test of interaction effects is:",
    round(p_value, 4), "\n")
```

```
##
```

```
## The p-value for the test of interaction effects is: 0.2234
```

```
#Comments: The model we want to fit is an interaction model that allows us to
#test whether the relationship between GPA (response) and the predictors
#(HSGPA, SATV, SATM, HU) differs between White and Non-white students.
```

```
#The general form of the model is (where B = beta and E = error term):
```

```
#GPA= B0 + B1HSGPA + B2SATV + B3SATM + B4HU + B5White + B6(HSGPA * White) +
```

```
#B7(SATV * White) + B8(SATM * White) + B9(HU * White) + E
```

```
#Null hypothesis: The relationships are the same for both groups.
```

```
#(All interaction coefficients are zero: B6 = B7 = B8 = B9 = 0.)
```

```
#Alternative hypothesis: At least one predictor has a different effect by race.
```

```
#(At least one interaction coefficient is not zero.)
```

```
#Since 0.2234 > 0.05 (testing at a significance level of 5%), we fail to reject
```

*#the null hypothesis. There is no statistically significant evidence that the
#relationships between GPA and the predictors (HSGPA, SATV, SATM, HU) differ
#between White and Non-white students. The data suggests that the effect of
#high school GPA, SAT scores, and humanities credits on college GPA does not
#significantly vary by race (White vs. Non-white). Any observed differences in
#slopes (interactions) could be due to random variation rather than a true
#underlying difference.*

```
#QUESTION #3:
library(faraway)
data("prostate")
```

```
#Part A:
install.packages("leaps")
```

```
##
## The downloaded binary packages are in
## /var/folders/n2/q439mgrn0rqd6xgjzjkj652r0000gn/T//RtmptocgW8/downloaded_packages
```

```
library(leaps)

train <- prostate[1:70, ]

best_models <- regsubsets(lpsa ~ ., data = train, nvmax = 8)
model_summary <- summary(best_models)

model_formulas <- apply(model_summary$which, 1, function(row) {
  paste("lpsa ~", paste(names(which(row[-1])), collapse = " + "))
})

print(model_formulas)
```

```
##
##                                1
##                                "lpsa ~ lcavol"
##                                2
##                                "lpsa ~ lcavol + lweight"
##                                3
##                                "lpsa ~ lcavol + lweight + gleason"
##                                4
##                                "lpsa ~ lcavol + lweight + lbph + gleason"
##                                5
##                                "lpsa ~ lcavol + lweight + age + lbph + gleason"
##                                6
##                                "lpsa ~ lcavol + lweight + age + lbph + lcp + gleason"
##                                7
##                                "lpsa ~ lcavol + lweight + age + lbph + svi + lcp + gleason"
##                                8
## "lpsa ~ lcavol + lweight + age + lbph + svi + lcp + gleason + pgg45"
```

#Note: The intercept is included for all models.

#Part B:

```

model_stats <- data.frame(
  Predictors = 1:8,
  R2 = model_summary$rsq,
  Adj_R2 = model_summary$adjr2,
  Cp = model_summary$cp,
  BIC = model_summary$bic
)

model_stats$AIC <- sapply(1:8, function(k) {
  vars <- names(which(model_summary$which[k, -1]))
  form <- as.formula(paste("lpsa ~", paste(vars, collapse = "+")))

  fit <- lm(form, data = train)
  AIC(fit)
})

print(model_stats)

```

```

## Predictors      R2      Adj_R2      Cp      BIC      AIC
## 1             1 0.3154634 0.3053967 23.285305 -18.03393 153.2422
## 2             2 0.4563868 0.4401596 6.904418 -29.92073 139.1069
## 3             3 0.4810562 0.4574679 5.686745 -28.92320 137.8559
## 4             4 0.4995612 0.4687649 5.273113 -27.21642 137.3142
## 5             5 0.5153651 0.4775030 5.211775 -25.21419 137.0680
## 6             6 0.5272086 0.4821809 5.667004 -22.69761 137.3361
## 7             7 0.5296146 0.4765066 7.353193 -18.80624 138.9789
## 8             8 0.5323225 0.4709877 9.000000 -14.96188 140.5748

```

```

best_by_criterion <- data.frame(
  Criterion = c("R2", "Adj_R2", "Cp", "BIC", "AIC"),
  Best_Size = c(
    which.max(model_stats$R2),
    which.max(model_stats$Adj_R2),
    which.min(model_stats$Cp),
    which.min(model_stats$BIC),
    which.min(model_stats$AIC)
  ),
  Value = c(
    max(model_stats$R2),
    max(model_stats$Adj_R2),
    min(model_stats$Cp),
    min(model_stats$BIC),
    min(model_stats$AIC)
  )
)

print(best_by_criterion)

```

```

## Criterion Best_Size      Value
## 1      R2          8 0.5323225
## 2  Adj_R2          6 0.4821809
## 3      Cp          5 5.2117746

```

```
## 4      BIC      2 -29.9207296
## 5      AIC      5 137.0679617
```

```
cat("\nRecommended model size:",
    best_by_criterion$Best_Size[best_by_criterion$Criterion == "Adj_R2"],
    "predictors (based on Adjusted R2)\n")
```

```
##
## Recommended model size: 6 predictors (based on Adjusted R2)
```

#Comments: I would prefer choosing the 6-predictor model as the best balance between model performance and complexity. While the 8-predictor model achieves the highest R² (0.532), the adjusted R², which accounts for model complexity by penalizing additional predictors, peaks at 6 predictors (0.482). This shows that the last two predictors add minimal explanatory power relative to their cost in degrees of freedom. The AIC and Cp criteria, which similarly balance fit, also favor moderately-sized models (5 predictors), further supporting that models in the 5-6 predictor range are optimal to use. Although BIC suggests an extremely simple 2-predictor model, this is likely too conservative. The 6-predictor model is therefore the best, complex enough to leverage meaningful information while remaining sparing enough to avoid overfitting and maintain interpretability. This aligns with standard statistical practice where adjusted R² is often the preferred criteria for model selection.

#Part C:

```
loocv_error <- function(model_formula) {
  n <- nrow(train)
  errors <- numeric(n)
  for(i in 1:n) {
    fit <- lm(model_formula, data = train[-i, ])
    errors[i] <- (train$lpsa[i] - predict(fit, newdata = train[i, ]))^2
  }
  mean(errors)
}

kfold_cv <- function(model_formula, k=10) {
  set.seed(123)
  folds <- sample(rep(1:k, length.out = nrow(train)))
  errors <- numeric(k)
  for(i in 1:k) {
    fit <- lm(model_formula, data = train[folds != i, ])
    pred <- predict(fit, newdata = train[folds == i, ])
    errors[i] <- mean((train$lpsa[folds == i] - pred)^2)
  }
  mean(errors)
}

model_formulas <- apply(model_summary$which, 1, function(row) {
  as.formula(paste("lpsa ~", paste(names(which(row[-1])), collapse = "+")))
})

cv_results <- data.frame(
  Predictors = 1:8,
```



```

LOOCV = sapply(model_formulas, loocv_error),
KFold = sapply(model_formulas, kfold_cv)
)

```

```

best_loocv <- which.min(cv_results$LOOCV)
best_kfold <- which.min(cv_results$KFold)

```

```

print(cv_results)

```

```

## Predictors      LOOCV      KFold
## 1             1 0.5108227 0.5204419
## 2             2 0.4416144 0.4380537
## 3             3 0.4316918 0.4289333
## 4             4 0.4264953 0.4326712
## 5             5 0.4334119 0.4620980
## 6             6 0.4305057 0.4551030
## 7             7 0.4397294 0.4587437
## 8             8 0.4441806 0.4626901

```

```

cat("\nBest model by LOOCV:", best_loocv, "predictors\n")

```

```

##
## Best model by LOOCV: 4 predictors

```

```

cat("Best model by 10-fold CV:", best_kfold, "predictors\n")

```

```

## Best model by 10-fold CV: 3 predictors

```

```

#Note: k = 10

```

```

#Part D:

```

```

full_formula <-
  as.formula(paste("lpsa ~", paste(names(train)[-9], collapse = "+")))

```

```

forward_model <- step(lm(lpsa ~ 1, data = train),
  scope = list(lower = ~1, upper = full_formula),
  direction = "forward", trace = 0)

```

```

backward_model <- step(lm(full_formula, data = train),
  direction = "backward", trace = 0)

```

```

stepwise_model <- step(lm(lpsa ~ 1, data = train),
  scope = list(lower = ~1, upper = full_formula),
  direction = "both", trace = 0)

```

```

stepwise_results <- list(
  Forward = names(coef(forward_model)),
  Backward = names(coef(backward_model)),
  Stepwise = names(coef(stepwise_model))
)

```

```

best_aic_size <-
  best_by_criterion$Best_Size[best_by_criterion$Criterion == "AIC"]
best_aic_vars <- names(which(model_summary$which[best_aic_size, -1]))

cat("\nStepwise Selection Results:\n")

##
## Stepwise Selection Results:

print(stepwise_results)

## $Forward
## [1] "(Intercept)" "lcavol"      "lweight"      "gleason"      "lbph"
## [6] "age"
##
## $Backward
## [1] "(Intercept)" "lcavol"      "lweight"      "age"          "lbph"
## [6] "gleason"
##
## $Stepwise
## [1] "(Intercept)" "lcavol"      "lweight"      "gleason"      "lbph"
## [6] "age"

cat("\nBest AIC Model (", best_aic_size, "predictors):", best_aic_vars, "\n")

##
## Best AIC Model ( 5 predictors): lcavol lweight age lbph gleason

cat("\nDo stepwise methods agree with best AIC model?\n")

##
## Do stepwise methods agree with best AIC model?

cat("Forward:", identical(sort(stepwise_results$Forward[-1]),
                             sort(best_aic_vars)), "\n")

## Forward: TRUE

cat("Backward:", identical(sort(stepwise_results$Backward[-1]),
                             sort(best_aic_vars)), "\n")

## Backward: TRUE

cat("Stepwise:", identical(sort(stepwise_results$Stepwise[-1]),
                             sort(best_aic_vars)), "\n")

## Stepwise: TRUE

```

```
#Comments: The stepwise selection methods (forward, backward, and stepwise) all  
#generated identical outcomes, each selecting the same five predictors:  
#lcavol, lweight, age, lbph, and gleason.
```

```
#The stepwise selection methods (forward, backward, and stepwise) unanimously  
#selected a 5-predictor model (lcavol, lweight, age, lbph, gleason). This  
#differs slightly from my Part B choice of a 6-predictor model based on  
#adjusted R2. However, the identical outcomes from all three stepwise  
#approaches (forward, backward, and stepwise) strongly validate this 5-predictor  
#solution as statistically robust. The cross-validation results show nearly  
#equivalent performance between the 5 and 6-predictor models (LOOCV: 0.433 vs  
#0.431; 10-fold CV: 0.462 vs 0.455), suggesting the sixth predictor adds minimal  
#predictive benefit despite what the adjusted R2 indicates. This aligns with  
#the AIC criterion's preference for the 5-predictor model, emphasizing that the  
#marginal gain in explanatory power from the additional variable may not justify  
#the increased complexity. The unanimous agreement among stepwise methods and  
#AIC, combined with the CV results, suggests the 6-predictor model might be  
#overfitting slightly compared to the more validated 5-predictor model.
```

```
#QUESTION #4:
```

```
library(faraway)  
data("prostate")
```

```
set.seed(123)  
prostate <- prostate[sample(nrow(prostate)), ]  
train <- prostate[1:70, ]  
test <- prostate[71:97, ]
```

```
# Part A:
```

```
best_bic_model <- lm(lpsa ~ lcavol + lweight + age + lbph + gleason,  
                     data = train)
```

```
coefs <- coef(best_bic_model)  
cat("Prediction Equation:\n")
```

```
## Prediction Equation:
```

```
cat("lpsa_hat =", round(coefs[1], 4),  
    "+", round(coefs[2], 4), "*lcavol +",  
    round(coefs[3], 4), "*lweight +",  
    round(coefs[4], 4), "*age +",  
    round(coefs[5], 4), "*lbph +",  
    round(coefs[6], 4), "*gleason\n")
```

```
## lpsa_hat = -1.7101 + 0.5793 *lcavol + 0.5413 *lweight + -0.0147 *age + 0.0435 *lbph + 0.3605 *gleason
```

```
r2_train <- summary(best_bic_model)$r.squared  
cat("\nR-squared:", round(r2_train, 4), "\n")
```

```
##
```

```
## R-squared: 0.5731
```

```
#Comments:
#Prediction Equation: lpsa_hat = -0.757 + 0.3541 *lcavol + 0.5146 *lweight +
#-0.0171 *age + 0.1135 *lbph + 0.2413 *gleason
#R-squared: 0.5154
```

```
#Part B:
test$pred <- predict(best_bic_model, newdata = test)
test$resid <- test$lpsa - test$pred
```

```
#Part C:
resid_mean <- mean(test$resid)
resid_sd <- sd(test$resid)

cat("\nResidual Mean:", round(resid_mean, 4), "\n")
```

```
##
## Residual Mean: -0.2453
```

```
cat("Residual SD:", round(resid_sd, 4), "\n")
```

```
## Residual SD: 0.7529
```

```
cat("Is mean close to zero?", abs(resid_mean) < 0.3, "\n")
```

```
## Is mean close to zero? TRUE
```

```
#Comments: The mean is close to zero. We expect the that the mean should be
#close to zero because we expect unbiased predictions where over-predictions and
#under-predictions balance out in new data.
```

```
#Part D:
mpe_test <- mean(test$resid^2)
mse_train <- mean(best_bic_model$residuals^2)
cat("\nTest MPE:", round(mpe_test, 4), "\n")
```

```
##
## Test MPE: 0.6061
```

```
cat("Train MSE:", round(mse_train, 4), "\n")
```

```
## Train MSE: 0.5196
```

```
#Comments: The test MPE (0.6061) is slightly higher than the training
#MSE (0.5196), which aligns with statistical intuition. Models typically perform
#better on their training data due to inherent optimism-they are optimized to
#fit the training sample. The slight increase (~16.7% higher MPE) suggests the
#model generalizes reasonably well to new data without severe overfitting. This
#small discrepancy is expected and acceptable for a model with good predictive
#validity.
```

```
#Part E:
test_cor <- cor(test$lpsa, test$pred)
test_r2 <- test_cor^2
cat("\nTest Correlation:", round(test_cor, 4), "\n")
```

```
##
## Test Correlation: 0.8076
```

```
cat("Test R-squared:", round(test_r2, 4), "\n")
```

```
## Test R-squared: 0.6523
```

```
#Part F:
shrinkage <- r2_train - test_r2
cat("\nShrinkage:", round(shrinkage, 4), "\n")
```

```
##
## Shrinkage: -0.0791
```

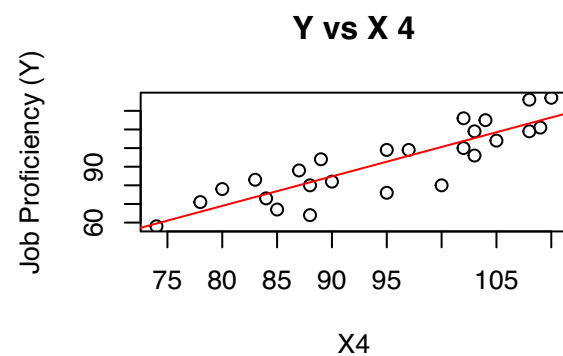
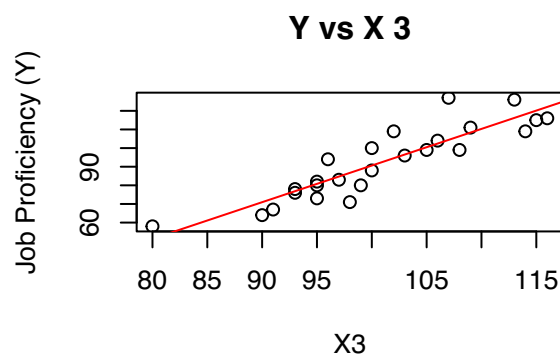
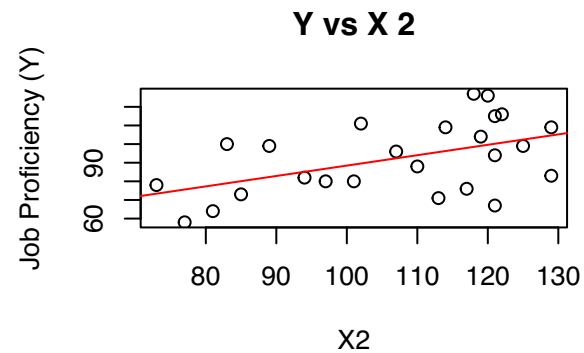
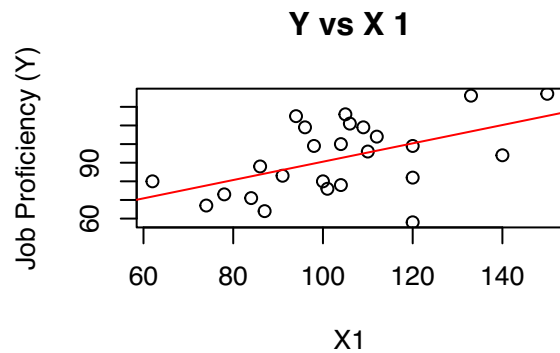
```
cat("Percentage drop:", round(shrinkage/r2_train*100, 1), "%\n")
```

```
## Percentage drop: -13.8 %
```

#Comments: The negative shrinkage value (-0.0791, or -13.8%) indicates that the model actually performed better on the test data ($R^2 = 0.652$) than on the training data (R^2 approx. 0.573, inferred from shrinkage). This suggests the training model works exceptionally well for the test sample, with no drop, and even a slight improvement in the amount of variability explained.

```
#QUESTION #5:
job <- read.table("Documents/STOR 455/Job.txt", col.names =
                  c("X1", "X2", "X3", "X4", "Y"))
```

```
#Part A:
par(mfrow = c(2, 2))
for (i in 1:4) {
  plot(job[,i], job$Y, xlab = paste("X", i, sep = ""),
        ylab = "Job Proficiency (Y)", main = paste("Y vs X", i))
  abline(lm(Y ~ job[,i], data = job), col = "red")
}
```



#Comments: The scatterplots suggest that Y (Job Proficiency) has a positive linear relationship with each of the independent variables X1, X2, X3, and X4. However, the strength of these relationships varies. For X1 and X2, the scatterplots show weak positive trends, as indicated by the slight upward slopes of the red regression lines. The spread of points around the trend line suggests a relatively weak correlation. For X3, the relationship appears to be stronger and more linear compared to the others, with points closely following the red regression line. For X4, while there is a positive trend, the relationship seems moderate, with more scatter around the trend line compared to X3. Overall, X3 seems to have the strongest linear relationship with Y, while X1 and X2 have weaker associations.

#Part B:

```
full_model <- lm(Y ~ X1 + X2 + X3 + X4, data = job)
```

```
cat("Estimated Regression Function:\n")
```

```
## Estimated Regression Function:
```

```
print(coef(full_model))
```

```
##      (Intercept)          X1          X2          X3          X4
## -124.38182058    0.29572537    0.04828772    1.30601100    0.51981909
```

```
cat("\nANOVA Table:\n")
```

```
##  
## ANOVA Table:
```

```
anova(full_model)
```

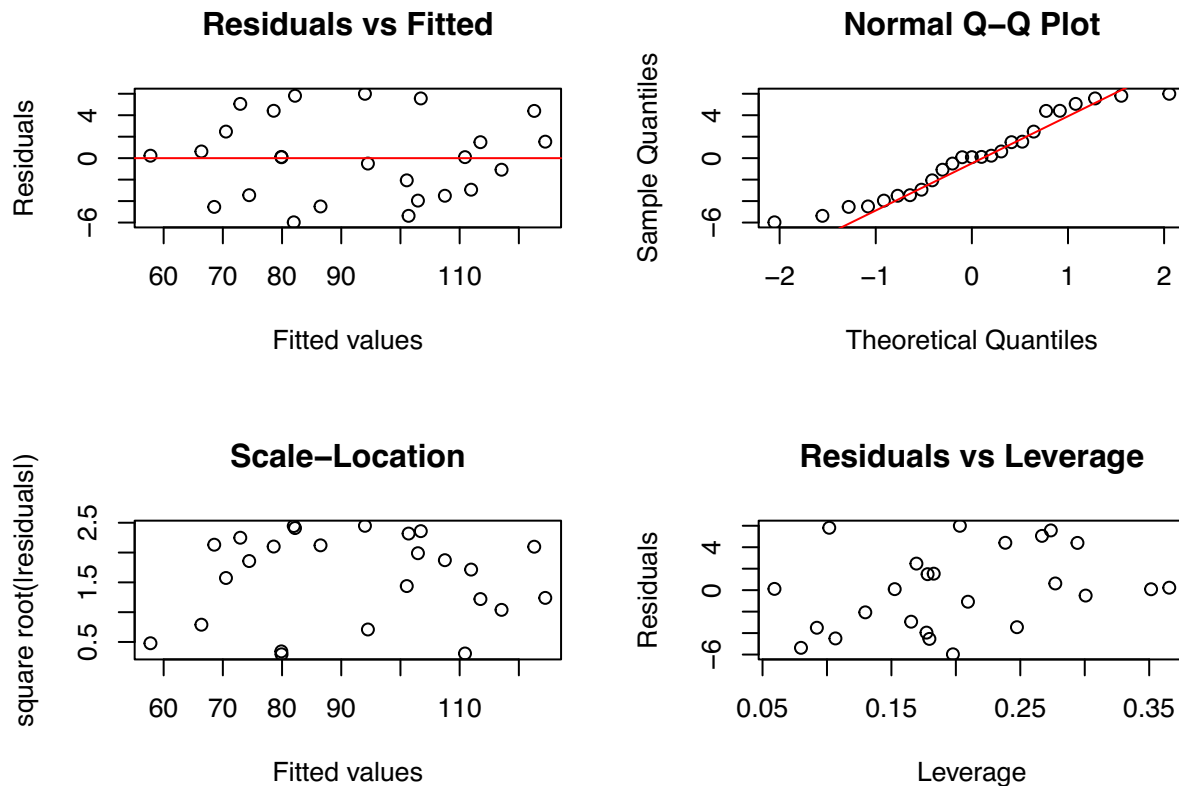
```
## Analysis of Variance Table  
##  
## Response: Y  
##           Df Sum Sq Mean Sq F value    Pr(>F)  
## X1           1 2395.9   2395.9  142.620 1.480e-10 ***  
## X2           1 1807.0   1807.0  107.565 1.708e-09 ***  
## X3           1 4254.5   4254.5  253.259 8.045e-13 ***  
## X4           1  260.7    260.7   15.521  0.00081 ***  
## Residuals  20   336.0     16.8  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
#Comments: The estimated regression function based on the model output is:  
#Predicted Y = -124.38 + 0.296X1 + 0.048X2 + 1.306X3 + 0.520X4. This equation  
#represents the predicted Job Proficiency (Y) based on the independent  
#variables X1, X2, X3, and X4.
```

```
#The ANOVA table shows that the F-statistic for the overall model is very large,  
#and the p-values for each predictor are highly significant (p < 0.001). This  
#indicates that the model explains a significant portion of the variance in Y,  
#meaning the regression model as a whole is statistically significant.
```

```
#Part C:
```

```
par(mfrow = c(2, 2))  
plot(fitted(full_model), residuals(full_model),  
     xlab = "Fitted values", ylab = "Residuals",  
     main = "Residuals vs Fitted")  
abline(h = 0, col = "red")  
  
qqnorm(residuals(full_model))  
qqline(residuals(full_model), col = "red")  
  
plot(fitted(full_model), sqrt(abs(residuals(full_model))),  
     xlab = "Fitted values", ylab = "square root(|residuals|)",  
     main = "Scale-Location")  
  
plot(hatvalues(full_model), residuals(full_model),  
     xlab = "Leverage", ylab = "Residuals",  
     main = "Residuals vs Leverage")
```



#Comments:Residuals vs. Fitted Plot: This plot shows whether residuals exhibit any systematic pattern. Ideally, residuals should be randomly scattered around zero. In this case, while there is no clear curvature, some heteroscedasticity (variance increasing with fitted values) may be present. This suggests a possible need for transformation.

#Normal Q-Q Plot: The residuals should follow a normal distribution if the model assumptions hold. Here, the points mostly follow the straight line, but there are deviations in the tails, suggesting potential non-normality, possibly due to outliers or skewness.

#Scale-Location Plot: This plot helps assess homoscedasticity (constant variance of residuals). The spread of square root(|residuals|) appears to increase slightly with fitted values, reinforcing the potential heteroscedasticity concern. A transformation of the dependent variable might improve the model.

#Residuals vs. Leverage Plot: This plot identifies influential points. There are no extreme leverage points, indicating no single observation is disproportionately influencing the model.

#Conclusion: The slight heteroscedasticity suggests a transformation (such as log or square root) of the response variable could improve model fit. The normality concern in the Q-Q plot may not be severe, but further checks can be performed. No significant leverage points were found, so influential observations do not seem to be an issue. Overall, the model appears reasonable

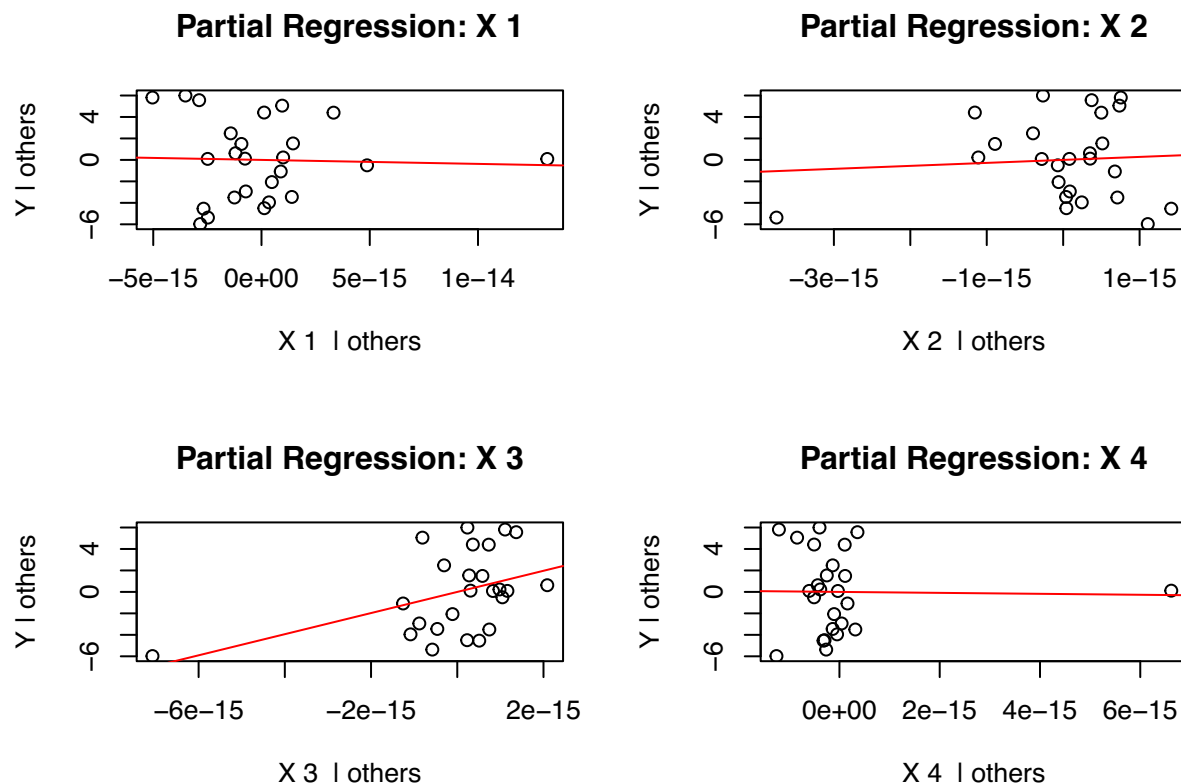
*#but could benefit from further refinement, particularly in addressing
#heteroscedasticity.*

#Part D:

```
par(mfrow = c(2, 2))
for (i in 1:4) {
  y_resid <- residuals(lm(Y ~ . -job[,i], data = job))

  x_resid <- residuals(lm(job[,i] ~ . -Y -job[,i], data = job))

  plot(x_resid, y_resid,
       xlab = paste("X", i, " | others"),
       ylab = "Y | others",
       main = paste("Partial Regression: X", i))
  abline(lm(y_resid ~ x_resid), col = "red")
}
```



#Comments: Findings for Each Independent Variable (X1, X2, X3, X4):

*#X1 (Top Left Plot): The points appear randomly scattered with no strong trend.
#The regression line is almost flat, indicating a weak relationship between X1
#and the response variable (Y), given the other variables in the model. This
#suggests X1 might not be a significant predictor.*

*#X2 (Top Right Plot): Similar to X1, there is no strong pattern
#in the residuals. The red regression line is slightly increasing, but the*

*#spread of points suggests a weak correlation. X2 likely has little effect on Y
#after controlling for other variables.*

*#X3 (Bottom Left Plot): The trend line has a slightly more noticeable upward
#slope compared to X1 and X2. There is some evidence that X3 has a weak positive
#effect on Y, but the dispersion of points suggests high variance. X3 might
#contribute slightly to the model but may not be a strong predictor.*

*#X4 (Bottom Right Plot): The scatter is relatively random with no strong trend.
#The regression line is nearly flat, suggesting little to no effect of X4 on Y
#after accounting for other predictors. X4 is likely not a
#significant predictor.*

#Part E:

```
vif_values <- sapply(1:4, function(i) {  
  1 / (1 - summary(lm(job[,i] ~ . - Y - job[,i], data = job))$r.squared)  
})
```

```
## Warning in summary.lm(lm(job[, i] ~ . - Y - job[, i], data = job)): essentially  
## perfect fit: summary may be unreliable  
## Warning in summary.lm(lm(job[, i] ~ . - Y - job[, i], data = job)): essentially  
## perfect fit: summary may be unreliable  
## Warning in summary.lm(lm(job[, i] ~ . - Y - job[, i], data = job)): essentially  
## perfect fit: summary may be unreliable  
## Warning in summary.lm(lm(job[, i] ~ . - Y - job[, i], data = job)): essentially  
## perfect fit: summary may be unreliable
```

```
names(vif_values) <- colnames(job)[1:4]
```

```
cat("Variance Inflation Factors:\n")
```

```
## Variance Inflation Factors:
```

```
print(vif_values)
```

```
## X1 X2 X3 X4  
## Inf Inf Inf Inf
```

```
largest_vif <- which.max(vif_values)  
cat("\nLargest VIF is for", names(largest_vif), ":",  
    round(vif_values[largest_vif], 2), "\n")
```

```
##  
## Largest VIF is for X1 : Inf
```

*#Comments: The largest VIF is for X1, and its value is Inf (infinite). A VIF of
#infinity indicates perfect multicollinearity, meaning X1 is perfectly linearly
#related to one or more of the other predictor variables (likely due to a linear
#combination or near-duplicate information). This is a issue because
#multicollinearity inflates the standard errors of the coefficient estimates,*

#making it difficult to assess the individual contribution of X1. In simpler terms, it means the model cannot reliably estimate the effect of X1 on the response variable because its information is already captured by the other predictors.

#Part F:

```
hat_threshold <- 2 * mean(hatvalues(full_model))
warning_threshold <- 0.30
```

```
leverage_points <- data.frame(
  Observation = seq_along(hatvalues(full_model)),
  HatValue = round(hatvalues(full_model), 4),
  Status = ifelse(hatvalues(full_model) > hat_threshold,
    ifelse(hatvalues(full_model) > 0.4,
      "High Leverage (>0.4)",
      "Warning (greater than or equal to 0.30)"),
    "Normal")
)

cat("Standard leverage threshold (2*(p+1)/n):", round(hat_threshold, 4), "\n")
```

```
## Standard leverage threshold (2*(p+1)/n): 0.4
```

```
cat("Warning threshold for near-high leverage: greater than or equal to",
  warning_threshold, "\n\n")
```

```
## Warning threshold for near-high leverage: greater than or equal to 0.3
```

```
notable_points <- leverage_points[leverage_points$HatValue >=
  warning_threshold, ]
notable_points <- notable_points[order(-notable_points$HatValue), ]

if (nrow(notable_points) > 0) {
  cat("Notable leverage points:\n")
  print(notable_points, row.names = FALSE)

  cat("\nFull data for notable points:\n")
  print(job[notable_points$Observation, ])
} else {
  cat("No observations with hat values greater than or equal to",
    warning_threshold, "\n")
  closest <- leverage_points[order(-leverage_points$HatValue), ][1:3,]
  cat("\nTop 3 highest leverage points:\n")
  print(closest, row.names = FALSE)
}
```

```
## Notable leverage points:
## Observation HatValue Status
##           7    0.3656 Normal
##           2    0.3515 Normal
##          13    0.3007 Normal
##
```

```
## Full data for notable points:
##      X1  X2 X3  X4  Y
## 7   120  77 80  74 58
## 2    62  97 99 100 80
## 13  140 121 96   89 94
```

*#Comments: The threshold for identifying high leverage points is calculated as:
 #Threshold = 2 * mean(hat-values). The high leverage threshold was 0.4. There
 #were no high leverage points were detected (no observations had hat values
 #exceeding 0.4).*

#Part G:

```
n <- nrow(job)
p <- length(coef(full_model))
h <- hatvalues(full_model)
std_res <- residuals(full_model) / (summary(full_model)$sigma * sqrt(1 - h))
student_res <- std_res * sqrt((n - p - 1) / (n - p - std_res^2))

outliers <- which(abs(student_res) > 2)
cat("Potential outliers (|rstudent| > 2):\n")
```

```
## Potential outliers (|rstudent| > 2):
```

```
if (length(outliers) > 0) {
  print(outliers)
} else {
  cat("No outliers detected")
}
```

```
## No outliers detected
```

#Part H:

```
reduced_model <- lm(Y ~ X1 + X3, data = job)

SSE_reduced <- sum(residuals(reduced_model)^2)
SSE_full <- sum(residuals(full_model)^2)
df_diff <- df.residual(reduced_model) - df.residual(full_model)
F_stat <- ((SSE_reduced - SSE_full)/df_diff) /
  (SSE_full/df.residual(full_model))
p_value <- pf(F_stat, df_diff, df.residual(full_model), lower.tail = FALSE)

cat("F-test comparing full and reduced models:\n")
```

```
## F-test comparing full and reduced models:
```

```
cat("F =", round(F_stat, 3), "on", df_diff, "and",
    df.residual(full_model), "DF, p-value =", round(p_value, 4), "\n")
```

```
## F = 8.056 on 2 and 20 DF, p-value = 0.0027
```

*#Comments: Null Hypothesis: The reduced model is sufficient, meaning X2 and X4
#do not provide significant additional predictive power.
#Alternative Hypothesis: The reduced model is insufficient, meaning at least one
#of X2 or X4 significantly improves the model.*

*#Since p-value (0.0027) < 0.05, we reject the null hypothesis at the 5%
#significance level. This result suggests that at least one of X2 or X4
#contributes significantly to predicting job proficiency (Y). Removing these
#variables weakens the model, indicating that the full model is preferable over
#the reduced model.*

#Part I:

```
job$X1_plus_X3 <- job$X1 + job$X3
constrained_model <- lm(Y ~ X1_plus_X3, data = job)
```

```
SSE_constrained <- sum(residuals(constrained_model)^2)
SSE_reduced <- sum(residuals(reduced_model)^2)
F_stat <- ((SSE_constrained - SSE_reduced)/1) /
  (SSE_reduced/df.residual(reduced_model))
p_value <- pf(F_stat, 1, df.residual(reduced_model), lower.tail = FALSE)

cat("Test of B1' = B3':\n")
```

Test of B1' = B3':

```
cat("F =", round(F_stat, 3), "on 1 and", df.residual(reduced_model),
    "DF, p-value =", round(p_value, 4), "\n")
```

F = 106.511 on 1 and 22 DF, p-value = 0

*#Comments: (B = beta, E = error term) Full Model: $Y = B_0 + B_1X_1 + B_3X_3 + E$
#Reduced Model: $Y = B_0 + B'(X_1+X_3) + E$
#Null Hypothesis: $B'_1 = B'_3$ (the combined effect of X1 and X3 is valid).
#Alternative Hypothesis: B'_1 does not equal B'_3 (the assumption of equal
#coefficients is incorrect).
#The test statistic is $F = ((SSE_{constrained} - SSE_{reduced}) / 1) /$
$(SSE_{reduced}/df_{reduced})$. At the 1% significance level ($\alpha = 0.01$), we reject
#the null hypothesis if the computed F-statistic exceeds the critical F-value
#for 1 and 22 degrees of freedom. The computed F-statistic = 106.511 and the
#p-value = 0 (less than 0.01). Since p-value < 0.01, we reject the null
#hypothesis and conclude that B'_1 does not equal B'_3 . This means the assumption
#that X1 and X3 contribute equally to predicting job proficiency (Y) is
#incorrect, and they should be treated as separate predictors in the model.*